

# The Use of Likely Invariants as Feedback for Fuzzers

Andrea Fioraldi<sup>1</sup>, Daniele Cono D'Elia<sup>2</sup>, Davide Balzarotti<sup>1</sup>

<sup>1</sup>EURECOM, <sup>2</sup>Sapienza University of Rome



[@andreafioraldi](https://twitter.com/andreafioraldi)

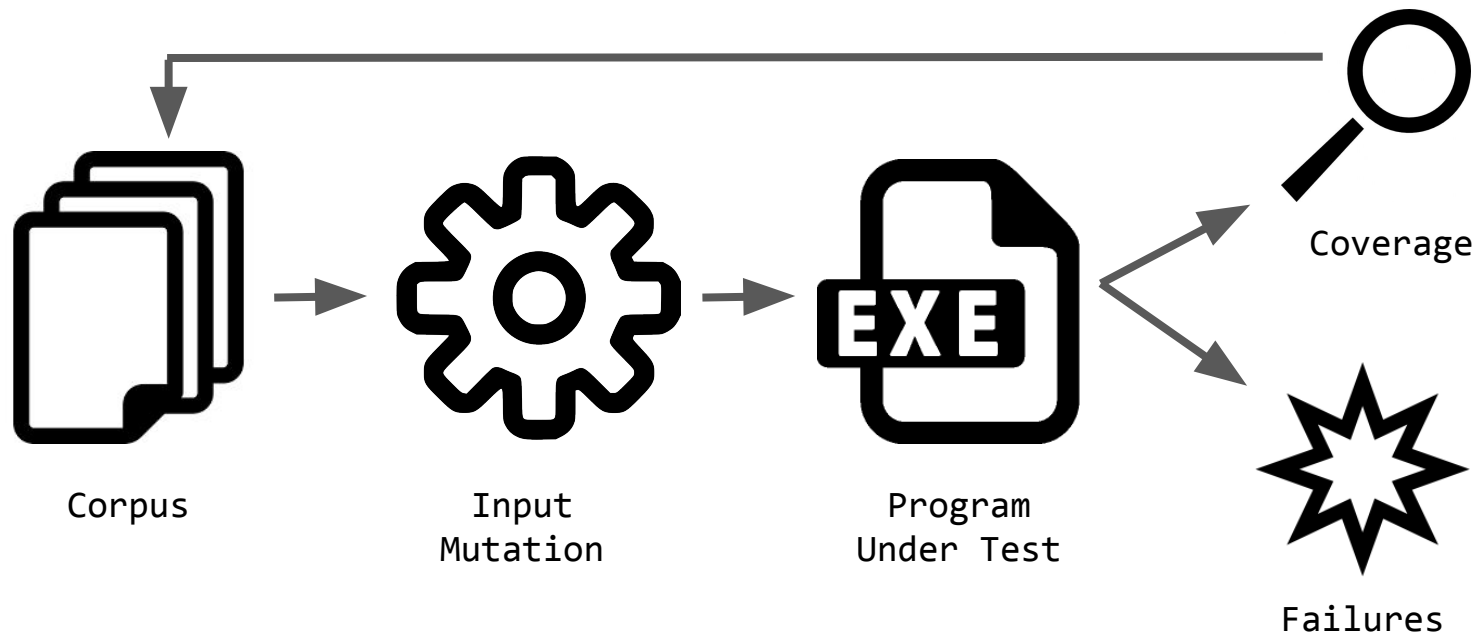


[fioraldi@eurecom.fr](mailto:fioraldi@eurecom.fr)



SAPIENZA  
UNIVERSITÀ DI ROMA


# Coverage-guided Fuzzing



# Coverage-guided Fuzzing

```
char buf[BUF_SIZE];
```


```
char foo(int x) {  
    if (x < sizeof(buf))  
        return buf[x];  
    return 0;  
}
```



# Coverage-guided Fuzzing

```
char buf[BUF_SIZE];
```


```
char foo(int x) {  
    if (x < sizeof(buf))  
        return buf[x];  
    return 0;  
}
```



# Coverage-guided Fuzzing

```
char buf[BUF_SIZE];
```

```
char foo(int x) {  
    if (x < sizeof(buf))  
        return buf[x];  
    return 0;  
}
```




# Coverage-guided Fuzzing

```
char buf[BUF_SIZE];
```

```
char foo(int x) {  
    if (x < sizeof(buf))  
        return buf[x];  
    return 0;  
}
```

Now the fuzzer cannot find any new coverage.



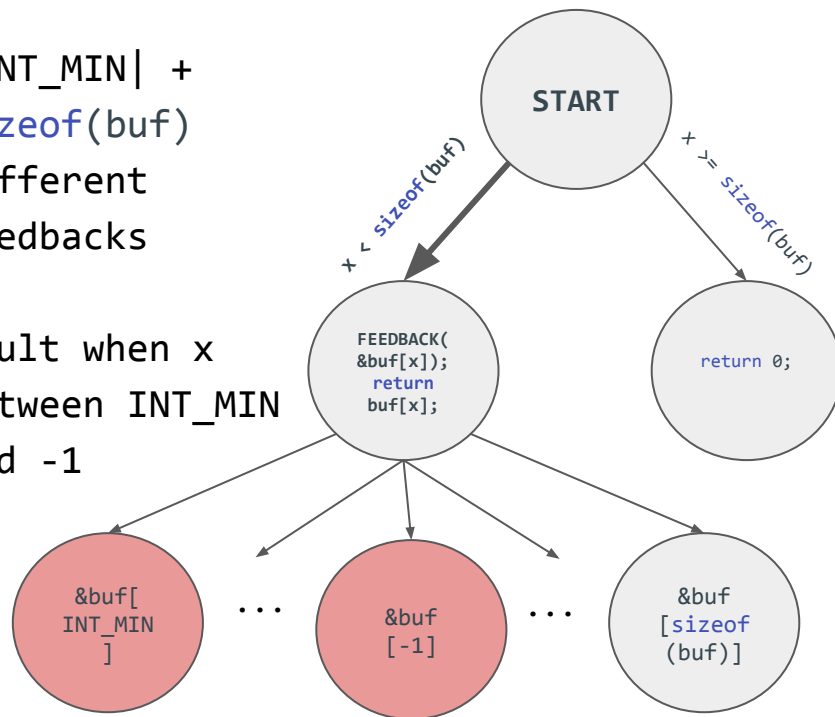
# Feedback-driven Fuzzing

```
char buf[BUF_SIZE];

char foo(int x) {
    if (x < sizeof(buf)) {
        FEEDBACK(&buf[x]);
        return buf[x];
    }
    return 0;
}
```

Note: x does not come directly from the input, otherwise it is easy for a fuzzer to trigger the bug thanks to the random generation

- `|INT_MIN| + sizeof(buf)` different feedbacks
- Fault when x between `INT_MIN` and `-1`




# Property-based Testing

```
char buf[BUF_SIZE];
```

```
char foo(int x) {  
    assert (x >= 0);  
    if (x < sizeof(buf))  
        return buf[x];  
    return 0;  
}
```

Note: `x` does not come directly from the input, otherwise it is easy for a fuzzer to trigger the bug thanks to the random generation





# Mining Invariants

- Learn invariants from static analysis
  - Over-approximation, no false positives, too generic invariants

# Mining Invariants

- Learn invariants from static analysis
  - Over-approximation, no false positives, too generic invariants
- Learn *likely* invariants from dynamic analysis
  - Needs a corpus of testcases
  - Suffers from the “*Coverage Problem*”, likely invariants are in worst case only local properties related to the observed executions
  - So many false positives

# Idea: Likely Invariants as Feedback


- In Fuzzing, we are interested in storing testcases that trigger new program states. Likely invariants violations can distinguish “unusual” states of the variables.

# Idea: Likely Invariants as Feedback

- In Fuzzing, we are interested in storing testcases that trigger new program states. Likely invariants violations can distinguish “unusual” states of the variables.
- Coverage Problem? Not a problem, local properties are interesting constraints to separate the space of the possible values.

# Idea: Likely Invariants as Feedback

- In Fuzzing, we are interested in storing testcases that trigger new program states. Likely invariants violations can distinguish “unusual” states of the variables.
- Coverage Problem? Not a problem, local properties are interesting constraints to separate the space of the possible values.
- Basic block, not function, level (ok for memory unsafe languages)



- $Y \in \{0, 2, 256\}$
- $Y < X$

The space is divided in 4 subspaces, A B C D

# Pruning Invariants

- Compute set of related variables and do not generate invariants between unrelated variables (*Comparability* calculation algorithm)

# Pruning Invariants

- Compute set of related variables and do not generate invariants between unrelated variables (*Comparability* calculation algorithm)
- Compute basic invariants using intra-procedural *Range Analysis* and instruct the learner to not output inviolable checks



# Pruning Invariants

- Compute set of related variables and do not generate invariants between unrelated variables (*Comparability* calculation algorithm)
- Compute basic invariants using intra-procedural *Range Analysis* and instruct the learner to not output inviolable checks
- Traverse the *Dominator Tree* in order to not emit the same check over the same variables different times in different basic blocks

# Prototype: InvsCov

Based on the AFL++ Fuzzing Framework and LLVM. We provide two LLVM passes and Clang wrappers.

Likely invariants mined using the Daikon dynamic invariants detector.



# Instrumentation

We extend the classic AFL-style instrumentation to keep track of the combination of violated invariants.

```
// Original AFL edge-coverage code
```

```
__afl_area_ptr[cur_loc ^ prev_loc]++;  
prev_loc = cur_loc >> 1;
```

```
// Extended to capture violations of invariants
```

```
__afl_area_ptr[cur_loc ^ prev_loc]++;  
prev_loc = cur_loc >> 1;  
prev_loc ^= __daikon_constr_1(var0);  
prev_loc ^= __daikon_constr_2(var2, var3);
```

```
...
```


# Evaluation

We evaluated our prototype using old versions of 8 real-world subjects used in previous works. The initial corpus is the result of a previous 24h fuzzing campaign, then we run the fuzzers for 48h.

- 62% growth on the fuzzer's queue size (no state explosion)
- 8% slowdown on the average testcase execution speed
- 1% growth on total edge coverage

# Evaluation

- 56% growth on bugs found vs. code coverage
- 41% growth on bugs found vs. context-sensitive coverage
- 67% of the bugs that only InvsCov could find were reached but not triggered by CodeCov



# Future Directions

- On-demand learning of invariants during fuzzer. Reimplement a fast version of Daikon that is also adaptive.

# Future Directions

- On-demand learning of invariants during fuzzer. Reimplement a fast version of Daikon that is also adaptive.
- Track implicit memory values not used in a block but still interesting (e.g. the size of buffers).

# Future Directions

- On-demand learning of invariants during fuzzer. Reimplement a fast version of Daikon that is also adaptive.
- Track implicit memory values not used in a block but still interesting (e.g. the size of buffers).
- Identify other ways to remove impossible-to-violate likely invariants at compile time.



# Future Directions

- On-demand learning of invariants during fuzzer. Reimplement a fast version of Daikon that is also adaptive.
- Track implicit memory values not used in a block but still interesting (e.g. the size of buffers).
- Identify other ways to remove impossible-to-violate likely invariants at compile time.
- What happens when we guide the fuzzer towards “rare” invariants violations?

Thank you!

Questions?